

Artifacts

- GitHub Repo:

<https://anonymous.4open.science/r/modis>

- Full Version:

https://anonymous.4open.science/w/modis_full/full.pdf

(From the 2nd page)

Multi-Objective Data Discovery for Data Science Models

ABSTRACT

Data discovery has been studied to identify relevant data to boost the performance of data science models in data-intensive tasks. In practice, there often comes the need to collect the data that can simultaneously optimize or trade among multiple performance metrics (such as training cost, accuracy, or inference cost) for specific models. Existing data augmentation techniques that apply table-wise data integration or feature selection are constrained to a single utility goal and meanwhile not sufficiently fine-grained to scale to multiple performance metrics. This paper introduces **MODis**, a novel paradigm to automatically create data from multiple sources with a goal to improve expected performance of a specific data science model in terms of multiple user-defined performance metrics. We formulate the problem as a multi-objective data discovery, and establish a formal computation system in terms of finite state transducer with finer grained operators enhanced by selection conditions. We analyze its expressiveness, and establish the hardness and approximability results. We then provide feasible algorithms that can create and maintain multiple datasets, from approximation scheme to fast heuristics that maintains the solution, all with guarantees in terms of Pareto optimality. We experimentally verify the efficiency and effectiveness of our algorithms in data creation. We also showcase their applications in accelerating the design of scientific data analytical pipelines.

1 INTRODUCTION

Data-driven analytical pipelines with data science models are routinely processed in a wide range of applications. Such pipelines rely on high-quality data science (machine learning) models. Among the challenges is the effective selection and creation of datasets that lead to high-quality models. In other words, *how to create new data to improve the overall (expected) performance of a model?*

Given a data science model M and a set of tables $\mathcal{D} = \{D_1, \dots, D_n\}$, data discovery aims to properly select *and* refine the data from \mathcal{D} to obtain a new data table D' , such that the expected performance of M over D' is improved. Moreover, the model M is evaluated by *multiple* (user-specified) performance measures such as training cost, processing cost, accuracy, inference cost, memory consumption, etc. It is desirable to suggest a dataset, over which the model has desirable performances in terms of *all* the measures.

Crowdsourced data platforms such as HuggingFace [14] provide portals to make datasets and models available. Data augmentation [30] and feature selection [22] have been also separately studied to improve machine learning, by carefully choosing useful data sources and feature space, respectively. Data integration has been adopted as an enabling technique to create new data for data augmentation [5]. Nevertheless, these approaches are not optimized where data science models are the first-class citizens. There is still a lack of effective solutions that (1) can suggest data that directly responds to a model as a “query” (i.e., “search data with a model”); and (2) in particular, improve the expected performances of the model in the presence of *multiple* performance measures.

We consider the following real-world example.

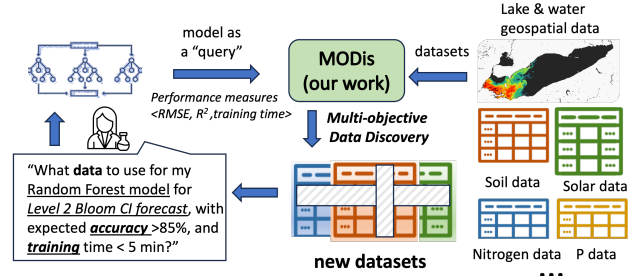


Figure 1: Discover and integrate features from multiple data sources improves the performance of a trained random forest (RF) for predicting CI index in both accuracy and training cost, as required by users. Manual tuning with cross-validation and feature selection tools find 10 features from more than 50 after several weeks. Our framework automatically suggested a new dataset with 8 features with comparable model accuracy and training cost, in a few minutes.

Example 1.1. To understand the impact of harmful algal bloom (HABs) of a lake and its reason, a research team needs a model that can forecast the chlorophyll-a index (CI-index), a critical algal bloom quantification index, in the short term. The team has collected a set of tables that involve more than 50 factors (e.g., fertilizer, water quality, weather) involving several upstream rivers and watershed systems of the lake. The team has already initialized a random forest (RF) model from small-scale tests, and wish to (1) improve its performance in terms of *multiple* metrics, including two accuracy measures: root mean square error (RMSE), and R^2 test, for “Level 2 bloom” CI-index, and its training cost; (2) understand *what* are new important factors and their values to be considered to achieve such improvement; and (3) track *where* the values are from and *how* they are integrated from the original data.

A typical process involves an excessive number of trials that repeats a pipeline of (manual) hyper-parameter tuning, feature selection, data integration, cross-validation, and inference tests. It turns out that only 8 key features with certain ranges of values are needed to achieve a regression with desirable accuracy at 89.6% (quantified by RMSE and R^2 test). This process took the team 8 weeks to get the desired model and the following observations.

- (1) Not all the water data tables from every upstream river system are needed, but only two - a basin area and an upstream river.
- (2) A single physicochemical attribute “Phosphorus” alone does not determine CI-index well for recent years, but “Nitrogen” and “Phosphorus” jointly determines CI-index, as an emerging pattern.
- (3) The original data contains, for both columns “Nitrogen” and “Phosphorus”, the values over the past 11 years. However, for an accurate short-term prediction in 2023, the fraction of the seasonal “Spring” data and in particular years “2013” and “2015” are important, due to high nutrition loads.

The above example calls for an automated data discovery process that can (1) suggest data from available ones to improve data science models for *multiple measures* at scale, and (2) provide an integrative

solution to suggest “what” data to choose and “how” to integrate the selected data. Moreover, the overall process should avoid expensive retraining and inference cost of the model.

One may exploit established data integration approaches to identify “joinable” tables and test the model. This requires an exhaustive enumeration of all the plans, leaving alone the overhead for actual evaluation. Such evaluation may already be expensive especially when M has high model complexity or with expensive user-defined functions [23], or involves nontrivial in-lab experiments or simulations as seen in scientific workflows [27]. Goal-driven data integration provides guided design for this process, yet may only respond to optimizing a single performance metric rather than multi-objective improvement for specific models.

Example 1.2. The research team may issue a query that asks: “What are proper datasets for which our random forest model can have an improved performance in forecasting ‘Level 2 bloom’ CI-index, with expected accuracy in terms of RMSE at most 0.3, R^2 score at least 0.7, and training cost less than 5 minutes?”

At first glance, this seems “contradicting” due to the trade-off between training cost and model accuracy for the measures on the CI index. Nevertheless, a data discovery process can still compromise to suggest one or more datasets, over which the model can “showcase” of desired accuracy for at least one metric (e.g., $a \geq 86\%$ accuracy for forecasting Level 2 bloom) and demonstrate comparable performance to the rest two metrics. This encourages us to seek for data discovery that can (1) *combine* data integration and multi-objective model evaluation, (2) pursue *multi-objective* optimization. Both are *new* challenges that are not well-addressed by existing data discovery methods for improving data-driven pipelines.

Moreover, the computation should not be limited to data augmentation by joining columns, but also to explore proper *selection conditions* and prune irrelevant rows. This reduces unnecessary overhead for downstream fine-tuning effort of the model.

This paper proposes a novel framework, from a formal characterization of the data discovery process to feasible algorithms with provable guarantees, to address the above two challenges. Our main technical contributions are as follows.

- (1) In response to the need of suggesting data with both columns and rows, We introduce a formal characterization of data discovery computation in terms of a *finite state transducer*, enhanced with fine-grained reduction and augmentation operators with selection conditions (Section 3). We study the expressiveness and property of the system. We show that the system resembles of data integration and feature selection. Better still, its computation demonstrate confluence property and a Church Russel property.
- (2) Based on our formulation of a data discovery system, we address the second challenge and introduce a novel problem of multi-objective data discovery problem (MODis) (Section 4). We model the problem as a multi-objective optimization process of the running of the data discovery system, which performs sequential operators to update an initial dataset. The running aims to outputs a *Pareto set* of datasets, over which a model is expected to have optimized performance over multiple evaluations metrics.
- (3) While the problem is in general intractable, we show that there

exists a fully polynomial time approximation scheme (FPTAS) for a constrained version of the problem (Section 5) under a fixed parameter specification. Our algorithm adopts a performance-driven search process, which interleaves multiple rounds of data integration and cost-effective model evaluation over historically observed tests and their performances.

Moreover, we introduce two variants of the algorithm to address the practical needs. The first exploits known correlations of the performance metrics to perform early termination and pruning, via a bi-directional search scheme (Section 6). The second introduce a diversification process to improve the generality of the generated datasets (Section 7). Both preserve the quality guarantee.

- (5) We finally introduce an efficient maintenance algorithm to maintain the datasets upon the updates of the underlying source dataset.

Using real benchmark datasets, models and analytical tasks, we experimentally verify the effectiveness of our data discovery scheme in improving the model performance and the efficiency of data generation. For example, our algorithms take 30 seconds to generate new data, that can improve pre-trained regression or classification models by 1.5-2 times in accuracy and makes their training cost 1.7 times faster. It outperforms other approaches that separately performs data augmentation or feature selection alone; and is feasible for large dataset. Our case study verifies its application in generating data for fine tuning data science models or benchmarking.

Related works. We categorize related works as follows. We remark that these work involve a common goal of improving the performance of data analysis models.

Feature Selection. Feature selection is a cornerstone task to remove irrelevant and redundant attributes and identify important ones for model training. Common strategies include filtering, wrapping, and embedding [22]. Filtering methods rank features in terms of correlation or mutual information [26, 29] and choose the top ones. They typically assume linear correlation among features, omitting collective effects from feature sets and hence are often limited to support directly optimizing model performance. Wrappers [17] choose features based on model performance estimated by a predictive model. The search cost is often extensive. Moreover, they risk overfitting due to optimizing a single criteria of model accuracy. Embedded methods directly learn feature selection models such as Lasso Regression, by penalizing (layer-wise) feature weights to improve model accuracy [36]. While they can be more efficient than wrapping, their performance are more sensitive to model architectures and task types, hyper-parameter tuning and additional domain-specific knowledge from datasets and models.

Our method differs in that (1) it creates new data with both data augmentation and masking of less useful values, rather than simply dropping the entire feature columns; (2) it finds data that improves the model in terms of multiple performance measures, beyond a single (accuracy) measure; and (3) the computation does not require internal knowledge of the models but simply data integration and masking operators, without incurring learning overhead. The process resembles SPJ (select, project, join) operations and readily benefit from well-established query optimization techniques.

Data Augmentation. Data augmentation and integration aim to create data from multiple data sources towards a unified view [5, 37].

It is often specified to improve data completeness and richness [30] and may be sensitive to the quality of schema that in turn requires entity matching [21] and schema matching [16]. Our method has a goal to create a dataset to improve the expected performance of given data-driven models and analytical pipelines. This is different from the goal of conventional data integration which mostly focuses on improving the completeness of data and query answers alone.

Closer to our work is goal-driven data discovery [8, 9]. The methods perform data augmentation queries (joins) to enrich input data to improve the performance of given tasks. The process discovers and augments datasets as path plans, and achieves the desired performance quality with a bounded number of queries. Our approach differs from these works as follows. (1) We formalize data discovery with a more expressive model with finer-grained column augmentation and cell-level operators with selection conditions. This allows us to create data with more expressiveness, and to better mitigate over-fitting or under-fitting issues. (2) We target multi-objective goals beyond a single performance measure. We provide algorithms with quality guarantees in terms of Pareto optimality, as well as optimization techniques. These are not addressed in prior methods.

Crowdsourced Data Services. Several platforms are available to allow users to share and search datasets, such as Dataset search [3], Data.gov [1], Kaggle [15], Hugging Face [14], and Zenodo [7]. These services exploit user-defined tags to retrieve relevant datasets, yet lack the necessary capability to provide datasets for improving the expected performance for specific models or tasks. Our approach is among the the first effort with the enhanced capability to directly create datasets that improve the expected performance of models as “queries”, and in terms of multiple performance measures.

2 MODELS AND PERFORMANCE EVALUATION

Datasets. A dataset $D(A_1, \dots, A_m)$ is a structured table instance that conforms to a relational schema $R_D(A_1, \dots, A_m)$. Each tuple $t \in D$ is a m -ary vector, where $t.A_i = a$ ($i \in [1, m]$) means the i th attribute A_i of t is assigned a value a . A dataset may have missing values at some attribute A (i.e., $t.A = \emptyset$).

We denote as $\mathcal{D} = \{D_1, \dots, D_n\}$ a set of given datasets. Each dataset D_i conforms to a schema R_i . We denote as \mathcal{A} the set of all the attributes from the datasets in \mathcal{D} . Given an attribute $A \in \mathcal{A}$, the *active domain* of A , denoted as $\text{adom}(A)$, refers to the finite set of all the values of A occurred in \mathcal{D} .

Models. We characterize a data science model as a function in the general form of $M : D \rightarrow \mathbb{R}^d$, which takes as input a dataset D , and outputs a result embedding in \mathbb{R}^d for some $d \in \mathbb{N}$. Here \mathbb{R} and \mathbb{N} are real and integer sets. In practice, M can be a pre-trained machine learning model, a statistical model, or a simulator. The input D may represent a feature matrix (a set of numerical feature vectors), or a tensor (from real-world physical systems), to be used for a data science model M as training or testing data. The output embedding can be conveniently converted to task-dependent output.

We say a model M is *fixed*, if its computation process does not change for fixed input. For example, a regression model M is fixed if any factors that determines its inference (e.g., number of layers, learned model weights) remain fixed. The model M is *deterministic* if it always outputs the same result for the same input. We consider

Symbol	Notation
\mathcal{D}, D, R_D	a set of datasets, dataset, and schema
$\mathcal{A}, A, \text{adom}(A)$	attribute set, attribute, and active domain
M	a data science model $D \rightarrow \mathbb{R}^d$
$\mathcal{P}, p, (p_l, p_u)$	performance measures, a measure, and its range
$T, t = (M, D, \mathcal{P}), t.\mathcal{P}$	a test set; a single test, and its performance vector
$\mathcal{T} = (s_M, \mathcal{S}, \mathcal{O}, \mathcal{S}_F, \delta)$	a data discovery system
\mathcal{E}	a performance estimation model
$C = (s_M, \mathcal{O}, M, T, \mathcal{E})$	a configuration of data discovery system
$G_{\mathcal{T}} = (\mathcal{V}, \delta)$	running graph
$s < s', D < D'$	state dominance, dataset dominance
\mathcal{D}_F	output (ϵ)-Pareto set

Table 1: Table of notations

fixed, deterministic data science models for the needs of consistency and accuracy required in data science pipelines.

Model Evaluation. Given a model M and an input dataset D , a *test* t is a triple (M, D, \mathcal{P}) that specifies a test dataset D , a pre-trained model M , and a set of (user-defined) *performance measures* $\mathcal{P} = \{p_1, \dots, p_l\}$. A *measure* p is an indicator to be *maximized* such as the model accuracy e.g., precision, recall, F1 (for classification); or *minimized* e.g., mean average error for regression tasks, training time, inference time, or memory consumption.

A performance measure $p \in \mathcal{P}$ can be always verified by an actual application of M over D , yet this can be expensive, especially when it involves in-lab tests or human inspection. We assume the measure p can be efficiently *approximated* by an estimation model \mathcal{E} (or simply “estimator”) [13, 35]. An estimator \mathcal{E} makes use of a set of historically observed performance vectors of M (denoted as T) to infer its performance over a new dataset. By default, we use a multi-output Gradient Boosting Model [28] that allows us to obtain the entire performance vector by a single call (see Section 8).

We use the following notations. (1) We unify \mathcal{P} as a set of normalized measures to be *minimized*, within a range $(0, 1]$. For a measure to be maximized (e.g., accuracy), one can easily convert it to an inversed counterpart (e.g., relative error). (2) Each measure $p \in \mathcal{P}$ has a user-specified pair of lower and upper bounds p_l, p_u to be normalized in $(0, 1]$, i.e., $0 < p_l \leq p_u \leq 1$. The range $[p_l, p_u]$ may specify tolerable training costs, memory consumption, or error ranges. (3) A test tuple $t = (M, D, \mathcal{P})$ is *valuated* by an estimator \mathcal{E} , which assigns a (estimated) value to each measure $p \in \mathcal{P}$ in polynomial time. The *performance vector* $t.\mathcal{P}$ is a valuated tuple where each entry has a constant $t.p$ of the corresponding metric $p \in \mathcal{P}$, estimated by \mathcal{E} .

Example 2.1. Consider Example 1.1. A pre-trained random forest (RF) model M that predicts CI-index is evaluated by three measures $\mathcal{P} = \{\text{RMSE}, R^2, T_{\text{train}}\}$, which specifies the root mean square error, the R^2 score, and the training cost. A user specifies a desired normalized range of RMSE to be within $(0, 0.6]$, R^2 in $[0, 0.35]$ w.r.t. a “inversed” lower bound $1 - 0.65$, and T_{train} in $(0, 0.5]$ w.r.t. an upper bound of “3600 seconds” (i.e., no more than 1800 seconds).

We summarize the main notations in Table 1.

3 DATA DISCOVERY: A FORMALIZATION

Given datasets \mathcal{D} and a fixed deterministic model M , we formalize a data discovery system in terms of finite state transducers. This formalization not only provides us the design principles of feasible

data discovery algorithms, but also verifies necessary properties that lead to their provable correctness and optimality guarantees.

3.1 Data Discovery System

We characterize a *data discovery system* as a finite state transducer, denoted as $\mathcal{T} = (s_M, \mathcal{S}, \mathcal{O}, \mathcal{S}_F, \delta)$, where

- \mathcal{S} is a set of states,
- $s_M \in \mathcal{S}$ is a designated start state,
- \mathcal{O} is a set of operators of types $\{\oplus, \ominus\}$;
- \mathcal{S}_F is a set of output states; and
- δ refers to a set of transitions.

We next specify its components.

States. A state s specifies a table D_s that conforms to schema R_s and active domains adom_s . For each attribute $A \in R_s$, $\text{adom}_s(A) \subseteq \text{adom}(A)$ refers to a fraction of values A can take at state s . $\text{adom}_s(A)$ can be set as empty set \emptyset , which indicates that the attribute A is not involved for training or testing M ; or a wildcard ‘_’ (‘don’t care’), which indicates that A can take any value in $\text{adom}(A)$.

Operators. We consider two classes of “shorthanded” operators to generate new data from an input dataset. These operators can be expressed and processed as SPJ (select, project, join) queries, enhanced with data augmentation or feature selection semantics.

(1) *Augmentation*: denoted as $\oplus_c(D_M, D)$, which augments a dataset D_M with dataset $D \in \mathcal{D}$ subject to a selection condition c . Here c is a single literal in the form of $A = a$ (an equality condition).

An augmentation operator $\oplus_c(D_M, D)$ performs the following queries: (a) augment the schema R_M of D_M with an attribute A from the schema R of D , if $A \notin R_M$; and (b) integrate D_M with tuples from D with values of $R.A$ satisfy value constraint c ; and (c) fill in the rest cells of the augmented table with “null” whenever their attribute values are not known.

(2) *Reduction* $\ominus_c(D_M)$: mask the cells of the tuples in D_M with the values of an attribute $R_M.A$ satisfy a literal constraint c with “null”. Here c is a single literal in the form of $A = a$, with $a \in \text{adom}(A)$.

Transitions. Given a state s and a set of operators \mathcal{O} of the two classes (simply denoted as $\{\oplus, \ominus\}$), applying an operator $op \in \mathcal{O}$ over s creates a new *result* state s' . We represent this as a *transition*, denoted as $r = (s, op, s')$, where s' is the result of applying op to s .

Running. A *configuration* of \mathcal{T} , denoted as $C = (s_M, \mathcal{O}, M, T, \mathcal{E})$, initializes a start state s_M with a dataset D_M , a finite set of operators \mathcal{O} of augmentation and reductions, a fixed deterministic model M , an estimator \mathcal{E} , and a test set T . Both D_M and T can be empty set \emptyset . A *running* of \mathcal{T} w.r.t. a configuration $C = (s_M, M, T, \mathcal{E})$ follows a general, deterministic process below.

- Starting from s_M , and at each state s , \mathcal{T} iteratively applies operators from \mathcal{O} , whenever applicable, to either augment a state with new attributes and tuples, or mask tuple values from the state. This spawns a set of children of a state, each specifies a new result dataset.
- For each transition $r = (s, op, s')$ that is spawned with a result s' by applying operator $op \in \mathcal{O}$, \mathcal{T} (1) constructs a test tuple $t(M, s'.D, \mathcal{P})$, (2) verifies if t is already in T ; if so, directly load $t.P$; otherwise, it invokes estimator \mathcal{E} at runtime to valuate t , and adds t to T .




Figure 2: Data Discovery Systems: Specifications

- The above process *terminates* at a set of output states \mathcal{S}_F , under the validation of an external termination condition, or no new transition can be spawned (no new datasets can be generated with \mathcal{O}).

The *result* of a running of \mathcal{T} refers to the set of corresponding datasets \mathcal{D}_F induced from the output states \mathcal{S}_F . As each output state $s \in \mathcal{S}$ uniquely determines a corresponding output dataset $s.D_s$, for simplicity, we shall use a single general term “output”, denoted as \mathcal{D}_F , to refer to output states or datasets as needed.

Example 3.1. Fig. 2 illustrates several classes of data discovery systems, categorized by the type of computation they perform.

(1) **Type 1** (resp. **Type 2**) system, denoted as \mathcal{T}_1 (resp. \mathcal{T}_2) starts from an initial dataset D_M in its start state s_M , and only applies reduction (resp. augmentation operators) to update D_M .

(3) **Type 3** system, denoted as \mathcal{T}_3 is a general form (that generalizes Types 1-2 systems) and perform either augmentation or reduction in any step to create datasets.

(4) **Type 4** system is a variant of Type 1 under a configuration that starts with a designated universal state s_U , which is associated with a “universal” dataset D_U . The dataset is obtained by concatenating all schema in \mathcal{D} as a “universal schema” R_U , and by integrating (joining) all tuples from \mathcal{D} . It can be specified to use a “reduce-from-universal” strategy and only performs reduction operators to remove cell values or tuples from D_U .

(2) **Type 5** system, denoted as \mathcal{T}_5 , extends transitions δ with an ϵ -transition (“do-nothing and jump”), so allows a “bi-directional” computation. It can perform (a) a reduction to remove values from an initial dataset followed by an option to (b) jump and start to augmentation to another dataset with new attributes and their values. For example, it may start with a universal dataset D_U as aforementioned and perform reduction; and choose to “jump” to augment a set of (small) datasets, or vice versa, until the computations converge to some same datasets.

Running graph. A running of \mathcal{T} can be naturally represented as the dynamic generation of a *running graph* $G_{\mathcal{T}} = (\mathcal{V}, \delta)$, which is a directed acyclic graph (DAG) with a set of state nodes \mathcal{V} , and a set of transition edges $r = (s, op, s')$. A *path* of length k is a sequence of k transitions $\rho = \{r_1, \dots, r_k\}$ such that for any $r_i = (s_i, op, s_{i+1})$, $r_{i+1} = (s_{i+1}, op, s_{i+2})$; i.e., it depicts a consecutive application of operators that converts an initial state s_1 with dataset D_1 to a new counterpart s_k with D_k . Consistently, we say a state node $s \in \mathcal{V}$ is *valuated*, if a corresponding test $t(M, D_s, \mathcal{P})$ is valuated by \mathcal{E} .

Example 3.2. Revisiting Example 1.1, the observations specify a small fraction of data that is sufficient for CI-index forecasting using a random forest (RF) M at desired accuracy. Fig 3 illustrates

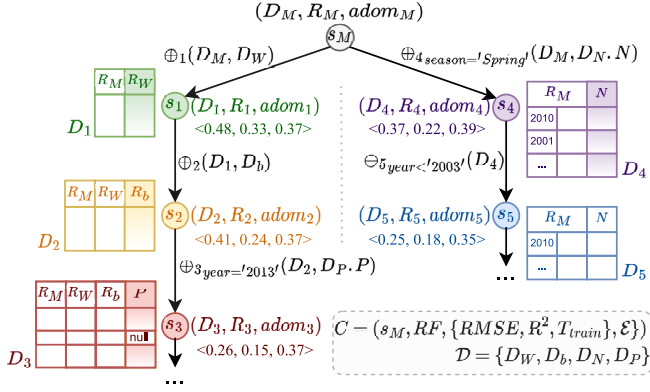


Figure 3: The running of a data discovery system, corresponding transition graphs, and result datasets.

a fraction of a running graph of a Type 4 system, creates such a dataset with \mathcal{D} contains four datasets $\{D_W, D_b, D_N, D_P\}$, for a water table, a basin table, a Nitrogen table, and a Phosphorus table. The augmentation \oplus is implemented with a library of spatial joins [31], a common query that join tables with tuple-level spatial similarity.

With a configuration $C = (s_M, RF, \{RMSE, R^2, T_{train}\}, \mathcal{E})$, where \mathcal{E} is an MO-GBM estimator \mathcal{E} , a running starts by joining D_W and D_b , the basin table and the upstream river data to get D_2 . It then augments D_2 with one more attribute “Phosphorus” under a selection condition “year = “2013”, and leads to result D_3 , via a path with edges labeled $\{\oplus_1, \dots, \oplus_3\}$. In each step, a test t is initialized (if not already in test set T); the estimator \mathcal{E} is consulted to evaluate the performance vector of t and enrich T .

Similarly, another path is spawned by the first converting D_M to D_4 with an augmentation, followed by a reduction that removes historical tuples older than “2003”, that are estimated to be less relevant. Fig 3 depicts a fraction of the transition graph of the above running, containing the two paths.

3.2 Expressiveness and Properties

We next discuss fundamental properties of a data discovery system. These properties help us justify the generality as well as practical implementations of data discovery algorithms (see Section 5).

Expressiveness. We study the expressiveness power of the data discovery system in terms of all the paths with operators as “labels” it can produce in all its possible transition graphs. This gives us a set of strings, forming its language $L(\mathcal{T})$. We show that a data discovery system and its runnings resembles, and can be specified in practice for data integration [20], or feature selection [24], We provide the following result.

PROPOSITION 3.3. *A data discovery system \mathcal{T} can be configured to express (1) data augmentation, and (2) feature selection.*

Proof sketch: We show the above cases by an analysis on the languages generated by the specifications of type 4 system \mathcal{T}_4 . To see this, it suffices to show that (1) $L(\mathcal{T}_1) \subseteq L(\mathcal{T}_4)$, and \mathcal{T}_1 can express feature selection; and (2) $L(\mathcal{T}_2) \subseteq L(\mathcal{T}_4)$, and \mathcal{T}_2 can simulate data augmentation process. One can readily infer (1) and (2) by treating \mathcal{T}_1 and \mathcal{T}_2 as special cases of \mathcal{T}_4 , with proper configurations. For (1), one removes the model M and constrain O as data reduction

operators only without selection conditions. For (2), one constraint O as data augmentation (“join”) operators. \square

Non-blocking. We also verify a desirable property of data discovery computation. A data operator (query) op is *non-blocking* [19], if for any path ρ with a transition that applies op and a result D in the running of a transducer \mathcal{T} , it generates the same result D regardless of at which step it applies op in ρ (whenever applicable). We have the following claim.

LEMMA 3.4. *The operators of type \oplus and \ominus are non-blocking.*

We observe that in most applications, (a) \oplus is commutative and associative, and (b) any consecutive application of an \oplus operator followed by an operator in \ominus are commutative. This ensures that the operators \oplus and \ominus are non-blocking. Moreover, this suggests that any paths in \mathcal{T} are order-independent.

Following the above analysis, we have the result below.

PROPOSITION 3.5. *Given a set of operators O with non-blocking operators \oplus and \ominus , Type 3, Type 4, and Type 5 systems have the same expressiveness. i.e., $L(\mathcal{T}_3) = L(\mathcal{T}_4) = L(\mathcal{T}_5)$.*

In other words, any dataset that can be created by a running of a general form \mathcal{T}_4 can be simulated by a running of \mathcal{T}_3 or \mathcal{T}_5 under a proper configuration that leads to the same result. This allows us to choose a proper design of \mathcal{T} from any of $\mathcal{T}_3, \mathcal{T}_4$ or \mathcal{T}_5 to perform the search with a completeness guarantee.

The above properties further indicate effective asynchronous implementation of data discovery in distributed environment for non-blocking operators. We defer such discussion in future work.

Church-Russel Property. With the above analysis, we verify that the computation of a data discovery system, from the perspective of a rewriting system, is both terminating and confluent, i.e., demonstrate a “Church-Russel” property.

PROPOSITION 3.6. *Given a configuration C with non-blocking operators of types $\{\oplus, \ominus\}$, and a data discovery system \mathcal{T} , the running of \mathcal{T} is terminating and confluent.*

We show this by verifying that: (1) any computation of \mathcal{T} leads to a “fixpoint” dataset D_F that has a universal schema R_U with all the tuples having ‘null’ values only, given possible and applicable operators in \oplus or \ominus . This is because augmentation does not repeatedly add new attributes from R_U , and each cell’s value, once set to be ‘null’ by a reduction, will not be changed again. (2) The sequences of operations are locally confluent (Lemma 3.4).

Due to limited space, we provide all the detailed proofs in [2].

4 MULTI-OBJECTIVE DATA DISCOVERY

Given a configuration, we want to find a running of \mathcal{T} that leads to a “global” optimal dataset, over which M is expected to achieve the best performance for all measures. Nevertheless, such an optimal solution may not always exist. (1) Even normalized to be minimized, the measures in \mathcal{P} may in nature conflict as trade-offs, such as training cost versus accuracy, or precision versus recall. (2) Given the “no free lunch” theorem [32], it is likely that no data-driven (ML) model performs better than others. We thus instead consider a scheme that pursues *Pareto optimality* for \mathcal{D}_F .

To this end, we introduce a state dominance relation below.

Dominance. Given \mathcal{P} and a data discovery system \mathcal{T} , a state $s = (D_s, R_s, \text{adom}_s)$ is *dominated* by another state $s' = (D'_s, R'_s, \text{adom}'_s)$ (resp. s' dominates s), denoted as $s < s'$ (resp. $s' > s$), if there are two valuated tests $t = (M, D_s)$ and $t' = (M, D'_s)$ in \mathcal{T} , such that

- for each measure $p \in \mathcal{P}$, $t.p \geq t'.p$; and
- there exists a measure $p^* \in \mathcal{P}$, such that $t.p^* > t'.p^*$.

Consistently, we say a dataset D_s is dominated by another D'_s , denoted as $D_s < D'_s$, if for their states s and s' , $s < s'$.

Given \mathcal{T} and a configuration C , let \mathcal{D}_F be the set of all the possible output datasets that can be generated by a running of \mathcal{T} , a set of dataset $\mathcal{D}_F^* \subseteq \mathcal{D}_F$ is a *Pareto set w.r.t. \mathcal{T} and C* , if

- any pair D_1 and D_2 in \mathcal{D}_F^* cannot dominate each other; and
- for any other $D \in \mathcal{D}_F \setminus \mathcal{D}_F^*$, and any $D' \in \mathcal{D}_F^*$, $D < D'$.

We now formulate the multi-objective data discovery problem.

Problem Statement. Given a configuration $C = (s_M, O, M, T, \mathcal{E})$, and a data discovery system \mathcal{T} that specifies non-blocking operators O and δ , the *multi-objective data discovery problem*, denoted as MODis, is to compute a Pareto set \mathcal{D}_F in terms of \mathcal{T} and C .

Example 4.1. Revisiting Example 3.2 and consider the temporal results $\mathcal{D}_F = \{D_1, \dots, D_5\}$ with the following performance vectors valuated by the estimator \mathcal{E} so far:

T: (D, M, \mathcal{P} , \mathcal{E})	RMSE	\hat{R}^2	T_{train}
$t_1 : (D_1, \text{RF}, \mathcal{P}, \text{MO} - \text{GBM})$	0.48	0.33	0.37
$t_2 : (D_2, \text{RF}, \mathcal{P}, \text{MO} - \text{GBM})$	0.41	0.24	0.37
$t_3 : (D_3, \text{RF}, \mathcal{P}, \text{MO} - \text{GBM})$	0.26	<u>0.15</u>	0.37
$t_4 : (D_4, \text{RF}, \mathcal{P}, \text{MO} - \text{GBM})$	0.37	0.22	0.39
$t_5 : (D_5, \text{RF}, \mathcal{P}, \text{MO} - \text{GBM})$	<u>0.25</u>	0.18	<u>0.35</u>

Here \hat{R}^2 is “inversed” as $1 - R^2$: the smaller, the better. All the measurement values are normalized in $(0, 1]$ w.r.t. user-specified upper and lower bounds, and the optimal values are underlined. One can verify the following dominance relation among the datasets: (1) $D_1 < D_2 < D_3$, and $D_4 < D_5$; (2) $D_3 \not< D_5$ and vice versa. Hence a Pareto set \mathcal{D}_F currently contains $\{D_3, D_5\}$, over the valuated cases.

The problem, in nature a multi-objective optimization problem, is in general intractable [10]. We next investigate feasible algorithms that implements the running of data discovery systems, with quality guarantees. We also investigate how different search strategies may impact the efficiency. Our main results are summarized in Table ??.

5 COMPUTING PARETO OPTIMAL DATASETS

To compute a Pareto set, a straightforward solution performs complete runnings of the system \mathcal{T} , to verify M 's performance over all the N states (datasets) and invoke established multi-objective optimization methods such as Kung's algorithm [18]. This will incur $O(N \log N)^{|\mathcal{P}|-2}$ number of valuation (for $|\mathcal{P}| \geq 4$), or $O(N(\log N))$ if $|\mathcal{P}| < 4$. Such valuation cost is infeasible in practice, even when enlisting N states as a “once-for-all” cost is affordable. Moreover, a Pareto set may already contain an excessive number of datasets that are too expensive to be inspected.

We next present an algorithm that approximates Pareto set with a size-bounded solution, and reduces unnecessary valuations.

5.1 Approximating Pareto Optimality

To characterize a feasible algorithm for data discovery with size-bounded solutions and quality guarantee, we introduce the notion of ϵ -Pareto set, followed by an optimality measure.

ϵ -Pareto set. Consider a data discovery system \mathcal{T} with a configuration (s_M, O, M) . Let \mathcal{D}_S be a set of datasets generated by a running of \mathcal{T} . Given two datasets D and D' in \mathcal{D}_S , and a constant $\epsilon > 0$, we say D ϵ -dominates D' , denoted as $D \leq_\epsilon D'$. If the corresponding tests $t=(M, D)$ and $t'=(M, D')$, $t.p \leq (1 + \epsilon)t'.p$ for each $p \in \mathcal{P}$.

A set of datasets $\mathcal{D}_\epsilon \subseteq \mathcal{D}_S$ is an ϵ -Pareto set of \mathcal{D}_S , if for every dataset $D' \in \mathcal{D}_S$, there exists a dataset $D \in \mathcal{D}_\epsilon$ such that $D \leq_\epsilon D'$.

(N, ϵ) -approximation. We say an algorithm is an (N, ϵ) -approximation for MODis, if it satisfies the following:

- for arbitrarily given, and fixed small constant $\epsilon > 0$, it correctly outputs an ϵ -Pareto set w.r.t. N valuated states; and
- the time cost is a polynomial determined by $|\mathcal{D}|$, N , and $\frac{1}{\epsilon}$.

Below we present our main result.

THEOREM 5.1. *Given datasets \mathcal{D} , configuration C , and a number N , there exists an (N, ϵ) -approximation for MODis in $O\left(\min(N_u^{|\mathcal{R}_u|}, N) \cdot \left(\left(\frac{\log(p_m)}{\epsilon}\right)^{|\mathcal{P}|-1} + 1\right)\right)$ time, where $|\mathcal{R}_u|$ is the total number of attributes from \mathcal{D} , $N_u = |\mathcal{R}_u| + |\text{adom}_m|$ with adom_m the largest active domain, $p_m = \max_{p \in \mathcal{P}} \frac{p_u}{p_l}$ as p ranges over \mathcal{P} ; and l the valuation cost per test.*

Remarks. The scheme characterizes a relative optimality w.r.t. ϵ and input size of states N . If N refers to the number of all possible states by an exhaustive running of \mathcal{T} , it ensures to output a ϵ -Pareto set of a Pareto set \mathcal{D}_F^* . Yet the algorithm is *not* in polynomial time w.r.t. dataset size $|\mathcal{D}|$ alone, as N can be exponential. On the other hand, it is still desirable as one can strike a balance between the ‘closeness’ of the output to a Pareto set, and the actual time cost, by explicitly tuning ϵ and a manageable size N .

5.2 Approximation Algorithm

As a constructive proof of Theorem 5.1, we next present an (N, ϵ) -approximation algorithm, denoted as ApxMODis.

“Reduce-from-Universal”. Algorithm ApxMODis implements a Type-5 system (Section 3.1) and starts with a node as a “universal” dataset D_U (with a universal schema). In a nutshell, it “transforms” state dominance to a “path dominance” counterpart, which aggregate valuated measures from nodes to the end of the paths at runtime. This strategy has the following desirable properties. (1) As we uniformly minimize all the measurements, this allows us to grow “shortest” paths with prioritized reductions and valuation to refine D_U towards any of the user provided lower bounds, encouraging *early termination*; (2) starting from more features and more tuples allow smaller sacrifice of model accuracy and generality.

Auxiliary structure. Algorithm ApxMODis dynamically spawns and maintains a running graph $G_{\mathcal{T}}$. In addition, it enhances $G_{\mathcal{T}}$ to a weighted graph, with the following global structure:

- (1) Each state node s is associated with a Boolean vector L (a bitvector) to encode if its local schema $s.R_s$ contains an attribute A from the universal schema, to guard the applicability of reduction, and to avoid to generate redundant dataset.

Algorithm 1 :ApxMODis

```

1: Input: Configuration  $C = (s_M, O, M, T, \mathcal{E})$ , a constant  $\epsilon > 0$ ,
   thresholds  $t$ ;
2: Output:  $\epsilon$ -Pareto set  $\mathcal{D}_F$ .
3: Queue  $Q := \emptyset$ ;  $Q.enqueue((s_M, 0))$ ;  $\mathcal{D}_F^0[pos(\rho_M)] = \mathcal{P}_{s_M}$ 
4: while  $Q \neq \emptyset$  do
5:    $(s', d) = Q.dequeue()$ ;  $\mathcal{D}_F^{d+1} = \mathcal{D}_F^d$ ;
6:   for all  $s \in \text{OpGen}(s')$  do
7:      $Q.enqueue((s, d + 1))$ ;
8:      $\mathcal{D}_F^{d+1} = \text{UPareto}(s, \mathcal{D}_F^{d+1}, \mathcal{D}_F^d, \epsilon)$ ;
9: return  $\mathcal{D}_F^{d+1}$ 
10: procedure OpGen( $s$ )
11:   set  $Q := \emptyset$ 
12:   for each  $l \in s.L$  do
13:     if  $l = 1$  then
14:        $s'.L = s.L$  with  $l$  set to 0
15:        $Q.append(s')$ 
16:   return  $Q$ 
17: procedure UPareto( $s, R, Q, \epsilon$ )
18:   for each  $\rho' \in Q$  do
19:     compute  $\mathcal{P}_s$  by invoking  $\mathcal{E}$ ; set  $\rho$  with  $\mathcal{P}_s$ ;
20:     for  $i \leftarrow 1$  to  $|\mathcal{P}_s|$  do
21:       if  $\mathcal{P}_s[i] > t[i]$  then continue;
22:     compute  $\text{pos}(\rho)$  with Equation (1);
23:     set  $p$  as deterministic metric
24:     if  $R[\text{pos}(\rho)] = \text{null}$  or  $R[\text{pos}(\rho)].p < \rho.p$  then
25:        $R[\text{pos}(\rho)] = \rho$ ;
26:   return  $R$ 

```

Figure 4: ApxMODis: Approximating Pareto Sets

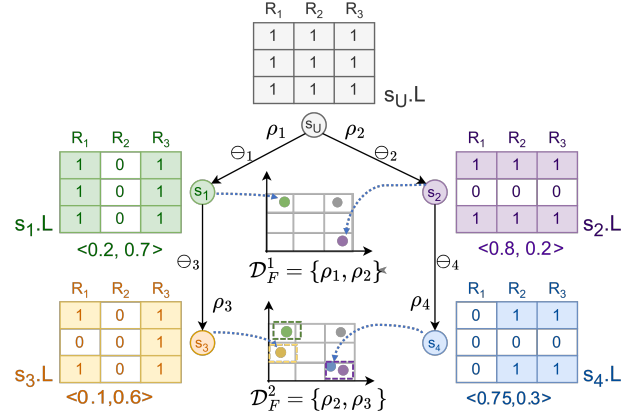
(2) Each transition (edge) $r = (s, \text{op}, s')$ has a vector of weights $\mathbf{c}(r)$ of length $|\mathcal{P}|$, where each entry $\mathbf{c}(r)_p = t.p - t'.p$ records the difference of valuated measure as p ranges over \mathcal{P} .

(3) As ApxMODis spawns paths, it also maintains at runtime for each path $\rho = \{r_1, \dots, r_k\}$ a *path label* as a pair $(\mathcal{P}(\rho), \text{pos}(\rho))$, where (a) $\mathcal{P}(\rho)$ is a $|\mathcal{P}|$ -ary vector with each entry $\mathcal{P}(\rho).p$ aggregates, for each measure $p \in \mathcal{P}$ as $\sum_{r \in \rho} \mathbf{c}(r).p$, and (b) $\text{pos}(\rho)$ “coordinates” ρ to a discretized $|\mathcal{P}|$ -ary space as

$$\text{pos}(\rho) = \left[\left\lceil \log_{1+\epsilon} \frac{\mathcal{P}(\rho).p_1}{p_{l1}} \right\rceil, \dots, \left\lceil \log_{1+\epsilon} \frac{\mathcal{P}(\rho).p_{|\mathcal{P}|}}{p_{l|\mathcal{P}|}} \right\rceil \right] \quad (1)$$

which captures a relaxed “envelope” of valuated performance, that allows us to decide ϵ -dominance by a simple comparison.

Algorithm. The algorithm ApxMODis is illustrated in Fig. 4, which induces a Pareto Set \mathcal{D}_F from a data discovery system \mathcal{T} . In *procedure* ApxMODis, we first construct a transition graph $G_{\mathcal{T}}$ using states \mathcal{S} and transitions δ , with a maximum depth of n . In line ?? to ??, for each state $s \in \mathcal{S}$, we initialize the initial layer of \mathcal{D}_{F_s} as an empty set and compute all $p \in \mathcal{P}$ for the current state s over M . At line ??, we assign the label of path ρ_M , with length 0, to position $\text{pos}(\rho_M)$ in $\mathcal{D}_{F_{s_M}}$. The iteration in line ?? to ?? is refining

**Figure 5:** “Reduct-from-Universal”: an illustration of two-level computation. The algorithm performs multiple level-wise spawns and updates ϵ -Pareto set.

the Pareto Set \mathcal{D}_F at each depth. When we reach a certain depth, for each state $s \in \mathcal{S}$, we carry over the Pareto Set from the previous round, then check all possible transitions that lead to s by Extend-&-Merge function, and update \mathcal{D}_{F_s} at the current depth accordingly. In *Extend-&-Merge function*, for each path ρ' that can potentially reach s , we generate $\rho = \rho' \cup \{s\}$ and evaluate it. If any of its measures in \mathcal{P} exceed the thresholds in \mathbf{ts} , we execute and proceed to the next path. Using Equation 1, we compute the ρ 's position in the Pareto set, which will be a $(d - 1)$ -vector without considering the deterministic metric. In line ?? to ??, we check the dominance with the paths at $\text{pos}(\rho)$ over all states by comparing the deterministic metric. This ensures that each position in the global Pareto set will only hold one most optimal path and keep the size of \mathcal{D}_F polynomially bounded. After iterating through all depths, the algorithm returns the refined Pareto set \mathcal{D}_F .

Example 5.2. Fig. 5 depicts a two rounds running of ApxMODis to create a $(5, 0.3)$ -Pareto set, for a given dataset $\mathcal{D} = \{D_1, \dots, D_3\}$ and a set of two measures $\mathcal{P} = \langle p_1, p_2 \rangle$. The operator set O contains four available reduction $\{\oplus_1, \dots, \oplus_4\}$. It starts with a start state s_U that specifies a universal dataset D_U with a universal schema R_U , which can be efficiently computed by optimized multi-way join. In the first round, it spawns a set of children from s_U that contains two states s_1 and s_2 , with two reductions Θ_1 and Θ_2 , respectively, as generated by the procedure OpGen by “flipping” the attributes and selection condition. It then consults the estimator \mathcal{E} to valuate their performances, and identified that $D_1 \not\leq_{0.5} D_2$ and vice versa. Thus it sets the current 0.3-Pareto set as $\{D_1, D_2\}$.

In the second round, it further spawns with applicable reductions Θ_3 and Θ_4 . By extending the path ρ_1 that leads to s_1 and ρ_2 that leads to s_2 , it obtains two new paths ρ_3 and ρ_4 with results D_3 and D_4 , respectively. The coordination finds that $D_3 \leq_{0.3} D_1$, but $D_2 \not\leq_{0.3} D_4$, and $D_2 \not\leq_{0.3} D_3$ and vice versa. This yields a new 0.5-Pareto set $\{D_2, D_3\}$. As the total valuated nodes reaches 5, it outputs $\{D_2, D_3\}$ as the result.

Correctness and Approximability. Algorithm ApxMODis terminates as it spawns over N nodes with at most $|R_U| |\text{adom}_m|$ distinct reduction operators, where $|\text{adom}_m|$ refers to the size of the largest

active domain. Following Proposition 3.5, it also correctly processes over a complete solution space that can be generated by the general Type-4 system with a Type-5 implementation.

We next verify the ϵ -approximability. For any constant ϵ , we show that our computation simulates an iterative process that solves a multi-objective shortest path problem (MOPS) [33]. Given an edge-weighted graph G_w , where each edge e_w is associated with a vector of utility values $e_w \cdot \mathcal{P}'$ and a start node u , it is to find a Pareto set Π of paths from u , where a path dominance relation between two paths ρ and ρ' is defined similar in terms of their utility vectors \mathcal{P}' , which is obtained by accumulating edge weight vectors. It has been shown that MOPS is fully polynomial time approximable (FPTAS) for the goal of computing ϵ -dominance. ApxMODis in nature simulates the FPTAS scheme with multiple rounds of “replacement” strategy in terms of path dominance relation.

Time cost. Let $|R_u|$ be the total number of attributes in the universal schema R_u of D_u , and $|\text{adom}_m|$ the size of the largest active domain. For time cost, ApxMODis performs $|R_u|$ levels of spawning, and at each node, spawns at most $|R_u| + |\text{adom}_m|$ children, given that it “flip” at most one attribute to reduce, and for each attribute, at most one domain value to mask. Let N_u be $|R_u| + |\text{adom}_m|$. Thus it valuates at most $\min(N_u^{|R_u|}, N)$ nodes (datasets), in $I \cdot \min(N_u^{|R_u|}, N)$ time, where I refers to a polynomial time valuation cost of \mathcal{E} per test. For each node, it then takes at most $\prod_{i=1}^{|\mathcal{P}|-1} \left(\left\lceil \log_{1+\epsilon} \frac{p_{u_i}}{p_{l_i}} \right\rceil + 1 \right)$ time to update the ϵ -Pareto set. Given ϵ is small, $\log(1 + \epsilon) \approx \epsilon$, and the total cost is in $O \left(\min(N_u^{|R_u|}, N) \cdot \left(\left(\frac{\log(p_m)}{\epsilon} \right)^{|\mathcal{P}|-1} + I \right) \right)$ time.

Given the above analysis, and consider $|R_u|$ and $|\mathcal{P}|$ are small constants, the cost is in a polynomial of input size $|D_u|$, N and $\frac{1}{\epsilon}$. Theorem 5.1 thus follows.

The analysis also verifies that ApxMODis provides a stronger guarantee when N is polynomially bounded by input size $|D_U|$ of the universal dataset. Given \mathcal{T} and configuration C , let \mathcal{D}_S be the finite set of all the datasets that can be generated by \mathcal{T} .

LEMMA 5.3. *If $|\mathcal{D}_S|$ is in $O(f(|D_U|))$ with f a polynomial, then the algorithm ApxMODis is a fully polynomial time approximation scheme (FPTAS) for MODis.*

We present the detailed analysis in [2].

6 “BI-DIRECTIONAL” DATA DISCOVERY

Algorithm ApxMODis may perform better when users specify higher expectations to the model performance (with smaller upper-bounds p_u), hence terminates earlier. For data discovery with larger $|\mathcal{D}|$, “reduction-only” ApxMODis may still valueate an excessive number of states. We next introduce a second algorithm, denoted as BiMODis, to perform a bi-directional search that implements Type-5 systems. In addition, it exploits a statistical correlation analysis of the “conflicting” measures to achieve early detection of dominance.

Algorithm. The algorithm, denoted as BiMODis, is illustrated in Fig. 6. Given a configuration, it performs the following.

Initialization (lines 3-4) It first initializes the following auxiliary structures: (1) two state nodes: a start state s_U associated with a

Algorithm 2 : BiMODis

```

1: Input: Configuration  $C = (s_U, \mathcal{O}, M, T, \mathcal{E})$ , a constant  $\epsilon > 0$ ;
2: Output:  $\epsilon$ -Pareto set  $\mathcal{D}_F$ .
3: queue  $Q_f := \{(s_U, 0)\}$ , queue  $Q_b := \{(s_b, 0)\}$ ,  $s_b = \text{BackSt}(s_U)$ ; /*output states of  $C^*$ */;
4: while  $Q_f \neq \emptyset, Q_b \neq \emptyset$  and  $Q_f \cap Q_b = \emptyset$  do
5:    $(s', d) = Q_f.\text{dequeue}()$ ; ▷ Forward Serach
6:    $(s'', d) = Q_b.\text{dequeue}()$ ; ▷ Backward Serach
7:    $\mathcal{D}_F^{d+1} = \mathcal{D}_F^d$ ;
8:   for all  $s^f \in \text{OpGen}(s')$  and all  $s^b \in \text{OpGen}(s'')$  do
9:     if  $\text{canPrune}(s^f, s^b)$  then
10:       Prune $(C, s^f, s^b)$ ;
11:        $\mathcal{D}_F^{d+1}.\text{add}(s^f)$ ;
12:        $\mathcal{D}_F^{d+1}.\text{add}(s^b)$ ;
13:    $\text{UPareto}(\mathcal{D}_F^{d+1}, \epsilon)$ ;
return  $\mathcal{D}_F$ ;

```

Figure 6: BiMODis: Bi-directional Search

universal dataset D_U , and by invoking a procedure BackSt, a backend state node s_b with a dataset D_b (see Procedure BackSt); (2) two queues Q_f and Q_b , to control the “forward” reduction of D_U , and “backward” augmentation of D_b . It also initializes the set \mathcal{D}_F to store the ϵ -Pareto set.

Bi-directional Search (lines 4-12). The bidirectional search is controlled by two queues Q_f and Q_b . In the forward stage, BiMODis applies reduction to extends paths and compute their coordinates as in ApxMODis. The backward search, on the other hand, applies augmentation to identify “promising” attributes to be augmented. This search is guided by a correlation analysis of features, which exploits a correlation network of features (not shown) to find highly relevant features, via Pearson correlation analysis (see details in ??).

Whenever a set of states are spawned and valuated from either forward stage or backward stage, it goes through a pruning check (by invoking a procedure canPrune; line 8) to early determine if an ϵ -dominance relation holds for any pairs of states, one from forward frontier, and the other from the backward. If so, it performs a pruning process to skip the spawning from one of the states. Otherwise, the bidirectional search continues to maintain \mathcal{D}_F .

The above process continues until no new states can be spawned from either the forward or the backward search, or at most N states are valuated. The result set \mathcal{D}_F is returned.

Procedure BackSt. This procedure initializes a backend dataset D_b to be augmented. This dataset can be task-specific. For example, for a classifier M with input features and a target attribute A to be classified, we sample a small (minimal) set of tuples in D_U to D_b , such that they cover all the values of adom of A , to ensure that no classes will be completely “missed” in D_b . Other strategies can be applied to initialize D_b .

We leave the details of the pruning strategy in [2].

Analysis. The algorithm BiMODis takes the same time cost as ApxMODis. The (N, ϵ) -approximation holds for BiMODis given that it correctly updates the ϵ -Pareto set by definition, and that it simulates the computation of Type 4 system. On the other hand, our

Algorithm 3 Procedure CorrFP (s, Rec, \mathcal{E})

```

1: Build  $G_c$  for measures recorded in  $Rec$ ;
2: StrongRs = GetSR( $G_c$ )
3: if  $s$  in  $Rec.keys()$  then                                ▶ Case 1: By  $Rec$ 
4:    $\mathcal{P}_s = Rec[s]$ ;
5:   if  $|valid(\mathcal{P}_s)| \geq 0.8|\mathcal{P}_s|$  then return  $\mathcal{P}_s$ ;
6: for all missing  $p_i^s$  in  $\mathcal{P}_s$  do                                ▶ Case 2: By  $G_c$ 
7:   if  $(p_i, p_j) \in \text{StrongRs}$  and  $p_j^s \in Rec[s]$  then
8:     find closed  $p_j^l$  and  $p_j^u$  with  $p_j^s$  in  $Rec$ ;
9:      $p_i^s = (p_i^l + p_i^u)/2$ 
10: if  $|valid(\mathcal{P}_s)| < 0.8|\mathcal{P}_s|$  then                                ▶ Case 3: By  $\mathcal{E}$ 
11:   fill missing  $p_s \in \mathcal{P}_s$  by invoking  $\mathcal{E}$ ; update  $Rec[s] = \mathcal{P}_s$ 
12: return  $\mathcal{P}_s$ ;

```

experimental study verifies that it is in practice much faster, and in particular suitable when creating new data from small datasets (see Section 8).

7 DIVERSIFIED DATA DISCOVERY

We have shown that feasible approximations exist for data discovery. In practice it is also desirable to diversify the datasets that are discovered, as recommending highly similar datasets may lead to biased training or testing data [11, 25].

We consider the following variant of MODis problem: Given a configuration C , a set of datasets \mathcal{D} , a number k , and constant ϵ , it is to compute an ϵ -Pareto set \mathcal{D}_F of at most k tables, that meanwhile maximizes a *diversification* goal as:

$$\max_{D, D' \in \mathcal{D}_F} \sum_{i=1}^k \sum_{j=i+1}^k d(D_i, D_j)$$

where function d computes a difference score between two tables.

We show that both algorithms ApxMODis and BiMODis can be readily extended to return a k -set of diversified ϵ -Pareto set, which remains to be (N, ϵ) -approximations. We outline such an algorithm, denoted as DivMODis is a variant of BiMODis.

Algorithm. The algorithm DivMODis solves a max-sum diversification problem on the fly while BiMODis is valuating a sequence of tables. It maintains a diversified subset $\mathcal{D}_F^P \subseteq \mathcal{D}_F$ with a fixed size k . It exploits a greedy strategy to gradually enlarge and update \mathcal{D}_F^P by deciding whether a valuated table from $\mathcal{D}_F \setminus \mathcal{D}_F^P$ at a level should join the diversified set or not.

Specifically, \mathcal{D}_F^P is initialized as \emptyset . Whenever $OpGen(s')$ (See line 8 in 6) is invoked, DivMODis updates \mathcal{D}_F^P with each newly spawned table D' in $OpGen(s')$, following two cases:

- (1) If $|\mathcal{D}_F^P| < k$, DivMODis adds $D' \in OpGen(s')$ that has the maximal marginal gain to \mathcal{D}_F^P ;
- (2) Otherwise, it scans each table in $OpGen(s')$ with a *sub-modular maximization solver* $\mathcal{S}_{\mathcal{A}}$ [4], which incrementally find an old table in \mathcal{D}_F^P to be replaced by D' . It is known that this greedy streaming selection strategy yields a $\frac{1}{4}$ -approximation for stream-based Max-Sum Diversification [6]. Differs from BiMODis, DivMODis updates

\mathcal{D}_F with the k -set \mathcal{D}_F^P to ensure it contains at least k diversified datasets, instead of feeding all spawned tables in $OpGen(s')$ to next level of bi-directional spawning. It thus returns a diversified ϵ -Pareto set \mathcal{D}_F w.r.t. N valuated.

Time cost. DivMODis also performs $|R_u|$ levels of spawning, and at each node, spawns and selects at most k children at each level. DivMODis incurs additional overhead at each level to update the diversified k -set. The cost of updating \mathcal{D}_F^P for a single level is in $O(k \cdot (|R_u| + |\text{adom}_m|) \cdot T_S)$ time, where T_S is a unit cost of invoking the solver $\mathcal{S}_{\mathcal{A}}$ to process a single table, which is in polynomial time. The total additional overhead for diversification of DivMODis is thus in $O(k(|R_u|^2 + |\text{adom}_m||R_u|T_S))$.

As k, R_u, adom_m and T_S are all small costs, the overhead for diversification is in practice small. As verified by our tests, DivMODis has comparable time cost with BiMODis in almost all cases (see Section 8). This verifies the feasibility of diversified data discovery.

8 EXPERIMENT STUDY

We next experimentally verify our algorithms. We aim to answer three questions: **RQ1**: How well can our algorithms improve the performance of models in multiple measures? **RQ2**: How fast they are in response to finding data for a model as “queries”? **RQ3**: What is the impact of configuration options? We also illustrate the applications of our approaches, with case studies.

Datasets. We use the following datasets summarized below:

- (1) Kaggle [15], a set of tables involving movie information;
- (2) OpenData: a fraction of a large open public dataset [1]. We sampled 2K tables, involving schools recording, school evaluations and school types, houses recording, housing price in New York and Chicago, geographical location, among others;
- (3) HF: a set of tables involving Avocado prices and relevant information, sampled from Hugging Face [14].

Dataset	Total #tables	Total #Columns	Total #Rows	Size
Kaggle	1943	33573	7317773	704.6MB
OpenData	2457	71416	33296998	8.58GB
HF	255	1395	10207261	3.5GB

Tasks and Models. We have trained the following models: (1) a random forest models RFhouse for classifying house price (Task T_1), using OpenData and the settings consistently in [8, 9]; (2) a Gradient Boosting Model (GBmovie) for predicting the movies’ worldwide gross sale, using Kaggle (Task T_2); and (3) a regression model for predicting Avocado price, using HF (Task T_3).

We trained all these models with scikit-learn [28]. For a fair comparison, we use the original training scripts provided by the baseline methods and validated that the reproduced models have consistent performance as reported.

Estimator \mathcal{E} . We trained a multi-output Gradient Boosting Model (MO-GBM) in scikit-learn [28] as estimator. This type of regressor outputs predicted values for multiple variables, thus allows us to valuate the entire performance vector for each test with one call.

Algorithms. We implemented the following algorithms in Python.

- (1) **MODis**: Our multi-objective data discovery algorithms, including ApxMODis, BiMODis, and DivMODis. We also implemented NOBiMODis, a counterpart of BiMODis without the optimization.

(2) **METAM** [9]: A goal-oriented data augmentation algorithm that perform consecutive joins of tables to optimize a single specific performance measure for a downstream task.

(3) **Starmie** [8]: A data discovery algorithm with table union search as the main use case. Given a table, it discovers union-able tables by learning feature semantic similarity with contrastive learning.

(4) **SkSFM** [28]: an automated feature selection method in scikit-learn's `SelectFromModel`, which recommend important features with a built-in estimator;

(5) **H2O** [12]: an AutoML platform used to automatically optimize machine learning pipelines. We utilized its feature selection module to fit features and predictors into a generalized linear model and then obtained a subset of selected features.

Evaluation. We evaluate data discovery algorithms in terms of tasks, performance measures and test models, as summarized below.

Tasks	Type	Dataset	Model	Perf.
T_1 : House Price (C)	Classification	OpenData	RFhouse	\mathcal{P}_1
T_2 : Movie Gross (C)	Classification	Kaggle	GBmovie	\mathcal{P}_2
T_3 : Avocado Price (R)	Regression	HF	LRavocado	\mathcal{P}_3

Performance Measures. For each task, we adopted a group of measures to guide the data discovery, as summarized below.

Notation	Measures	Used In
p_{Acc}	Model Accuracy	\mathcal{P}_1 - \mathcal{P}_3
p_{Tr}	Training Time Cost	\mathcal{P}_1 - \mathcal{P}_3
p_{F1}	F_1 score	\mathcal{P}_2
p_{MAE}	Mean Absolute Error	\mathcal{P}_3
p_{MSE}	Mean Squared Error	\mathcal{P}_3
p_{Fsc}	Fisher Score [22]	\mathcal{P}_1
p_{MI}	Mutual Information [9, 22]	\mathcal{P}_1
p_{VIF}	Variance Influence Factor [22]	\mathcal{P}_1
p_{DSize}	Size of Created Dataset	None

(1) The first four directly quantify a model's performance in terms of accuracy (p_{Acc} for classification and regression, and both p_{MAE} and p_{MSE} for regression) and training cost (p_{Tr}). (2) The latter three (p_{Fsc} , p_{MI} and p_{VIF}), generally used in feature selection [22] quantify the statistical relationship between a set of input variables (features) and a "target" feature (e.g., 'House Price' to be classified, or 'Avocado Price' to be predicted); the larger, the better. Among these, p_{MI} is also adopted by [9] as an optimization goal for data discovery. (3) To evaluate the amount of result, we also report the size of the data (p_{DSize}), in terms of (total # of rows, total # of features). As all baselines only report a single table, and MODis report a set of tables, we report total size in favor of baselines. Here if a column has all cell masked, we consider the column reduced and remove it from the output table.

For each tasks in T_1 - T_3 , we initialized our MODis methods consistently with a configuration that specifies an original dataset, the matching trained model, and the corresponding measures \mathcal{P}_1 - \mathcal{P}_3 .

Evaluation metrics. We adopt the following metrics to quantify the effectiveness of data discovery approaches. Denote as D_M an initial dataset, and D_o a set of output datasets from a data discovery algorithm. (1) We define the *relative improvement* $rlmp(p)$ for a given measure p achieved by a method as $\frac{M(D_M) \cdot p}{M(D_o) \cdot p}$. As all metrics are normalized to be minimized, the larger $rlmp(p)$ is, the better D_p is in improving M w.r.t. p . Here $M(D_M) \cdot p$ and $M(D_p) \cdot p$ are

obtained by actual model inference test. This allows us to fairly compare all methods in terms of the quality of data suggestion.

For efficiency, we compare the total time cost of the data discovery process upon receiving a given model or task as a "query".

Exp-1: Effectiveness with Single Measure. Our first experiment evaluates all algorithms in evaluating how well the model's performance can be improved over the dataset(s) they created. As p_{Acc} is the single measure considered by METAM, and all baseline produce a single table, we (1) compare MODis algorithms by selecting the table in the Pareto set with best estimated p_{Acc} , and (2) apply model inference to all the datasets, to report the actual measurement values. We show the results for T_3 in Table 2 ("Original" refers to the measures over the input dataset). We find the following.

(1) MODis algorithms outperform all the baselines in creating a dataset to improve the performance in terms of p_{Acc} . The one with best p_{Acc} and second best is obtained by BiMODis and NOBiMODis, respectively, and all MODis methods finds data for which p_{Acc} achieves 0.94.

(2) Over the same dataset and for other measures, MODis algorithms still outperforms the baselines in most cases. For example, the result datasets that optimize p_{Fsc} , p_{MI} and p_{VIF} are obtained by ApxMODis, NOBiMODis and DivMODis, respectively; and BiMODis finds a dataset that achieves three second best results in p_{Train} , p_{Acc} and p_{Fsc} . This verifies their ability in optimize data discovery towards multiple measures simultaneously.

(3) All baseline methods perform data augmentation or feature selection that leads to a single table. The data augmentation methods (METAM, Starmie) mainly include more features to improve accuracy; and feature selection (SkSFM and H2O) reduce them at a cost of accuracy but improved training cost. MODis methods are able to balance these trade-offs better by *explicitly* performing multi-objective optimization. Consider p_{Acc} and p_{Train} . The best result for training cost is contributed from SkSFM, yet at a cost of lowest model accuracy. As MODis methods are able to optimize both measures (among others), by making flexible decision to augment with new features or reduce cells and tuples to make the data smaller, they are able to find data with improved accuracy as well as smaller training cost, compared with baselines.

(4) Despite p_{Acc} is a first-class citizen in this comparison, not all baselines improve it (given its value over "Origin") significantly, except Starmie. Yet Starmie improves accuracy at a cost of including the most number of features (13 new ones). Feature selection methods (SkSFM and H2O) achieved better result on accuracy with much less number of features, and consistently showing better results in feature correlation measures in terms of p_{Fsc} , p_{MI} and p_{VIF} . On the other hand, MODis methods *explicitly* included these into optimization scope with a multi-objective estimator, and are able to improve accuracy without introducing many new attributes.

Exp-2: Effectiveness with multiple measures. We next evaluate MODis algorithms, METAM and Starmie, using multiple measures in T_1 and T_3 . For each measure p and an algorithm, we choose the dataset D with the best estimated measure of p it generates. We then retrain the model using D to get the true measurement. We normalize all the values into a same range. The results are illustrated

T_3 : Movie Gross (C)	Original	METAM	Starmie	SkSFM	H2O	ApxMODis	NOBiMODis	BiMODis	DivMODis
p_{Acc}	0.8560	0.8515	0.8606	0.8285	0.8545	0.9257	0.9782	<u>0.9755</u>	0.9413
p_{Train}	1.4775	1.5905	1.2643	0.6028	0.9692	1.0575	0.9107	<u>0.8027</u>	1.0486
p_{Fsc}	0.0824	0.0949	0.1286	0.7392	0.3110	0.6013	0.9257	<u>0.9240</u>	0.6618
p_{MI}	0.0538	0.0580	0.1072	0.3921	0.1759	0.4108	<u>0.3944</u>	0.3839	0.3644
p_{VIF}	1.5831	1.5831	<u>1.2980</u>	1.6742	1.5096	2.2023	1.9743	1.7688	1.2923
Output Data Size	(3264, 10)	(3264, 12)	(3264, 23)	(3264, 3)	(3264, 8)	(2958, 9)	(1935, 11)	(1835, 11)	(2894, 9)

Table 2: Comparison of Data Discovery Algorithms in Multi-Objective Setting

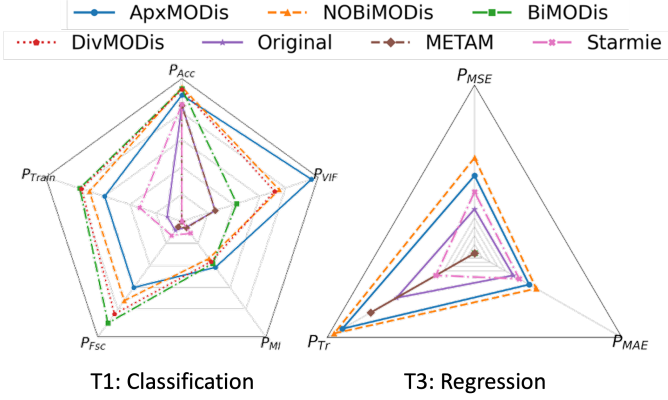


Figure 7: Effectiveness: Multiple Measures

as radar graphs in Fig. 7. The lines “Original” mark the values of the measures in the original data.

In general, MODis algorithms are able to create datasets that generally improve a model in a balanced performance. In particular, BiMODis, DivMODis and NOBiMODis provide top results for multiple measures. METAM is optimized to provide good results for a single measure, such as accuracy in T_1 . Starmie is not specifically optimized for optimizing measures, and provides a balanced performance in T_2 . ApxMODis provides in particular better results over p_{VIF} , a measure for feature correlation, with a possible reason that it performs more localized reduction only operations that is closer to feature selection process.

Exp-3: Effectiveness and Impact factors. We next investigate the performance of MODis methods under the impact of factors, in particular ϵ , the approximation factor; and the maximum path length (maxl) to bound the number of valuated states N .

Accuracy vs. ϵ . Fixing maxl = 6, we varied ϵ from 0.5 to 0.1. As shown in Fig. 8(a), MODis algorithms are able to improve the model in p_{acc} better with smaller ϵ , as they all ensure to output a ϵ -Pareto set that better approximate a Pareto set when ϵ is set to be smaller. In all cases, they achieve a relative improvement $rlmp(p_{Acc})$ at least 1.07. BiMODis and NOBiMODis perform better in recognizing better solutions from both ends in reduction and augmentation as smaller ϵ is enforced. ApxMODis, with reduction only, is less sensitive to the change of ϵ due to that larger ϵ may “trap” it to local optimal sets from one end. Adding diversification (DivMODis) is able to strike a balance between ApxMODis and BiMODis by enforcing to choose difference datasets out of local optimal sets, thus improving ApxMODis for smaller ϵ .

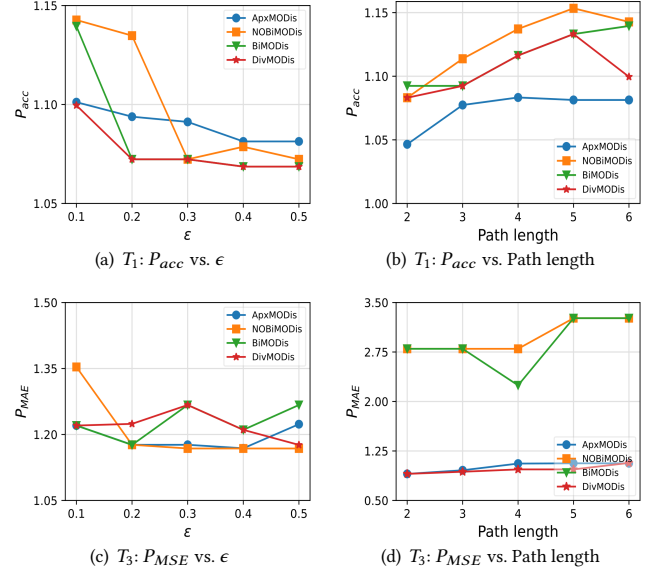


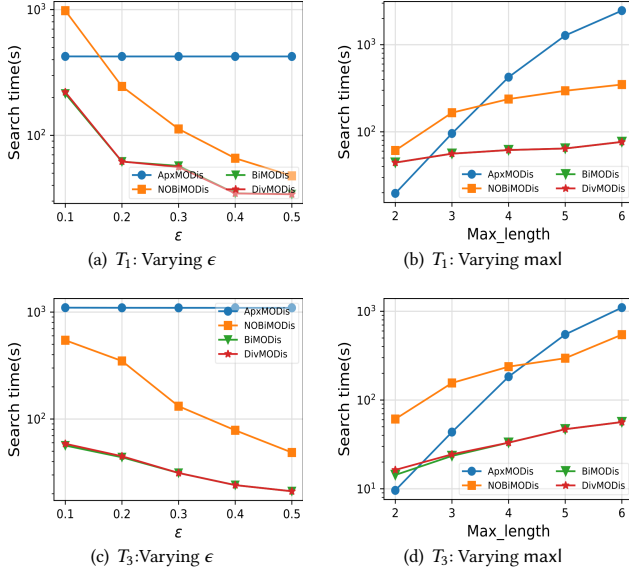
Figure 8: Effectiveness with Varying Factors

Accuracy vs. maxl. Fixing $\epsilon = 0.1$, we varied maxl from 2 to 6. Fig. 8(b) tells us that all MODis algorithms improves the accuracy of classification better for more rounds of path extension, as expected. Specifically, BiMODis and NOBiMODis benefit most as bi-directional search allow both to find better solution from wider search space as maxl become larger. ApxMODis and DivMODis are less sensitive, yet may due to different reason. ApxMODis is able to be robust in reporting dataset with similar quality due to localized search with reduction; and DivMODis, while can find the dataset with best result (e.g., when maxl = 5), may “lose chance” to keep the one by replacing it with less optimal but more different counterparts when its search enlarges (when maxl = 6).

Using HF data, we report the impact of ϵ and maxl for Task 4 for regression analysis in Figs. 8(c) and 8(d). The results are consistent with their counterparts in Figs. 8(a) and 8(b). As ϵ varies from 0.1 to 0.5 (with maxl fixed as 6), BiMODis is able to generate datasets that improve p_{mse} of the regression by 1.5 times.

Exp-3: Efficiency. We next report the efficiency of the MODis on generating datasets for task T_1 and task T_3 over OpenData and HF, respectively, and the impact of two major factors ϵ and maxl.

T_1 : Varying ϵ . Fixing maxl = 4 and varying ϵ from 0.1 to 0.5, Fig. 9(a) verify the following. (1) BiMODis, NOBiMODis and DivMODis

Figure 9: Efficiency with Varying Factors (T_1)

take less time to create datasets as ϵ becomes larger. As ϵ is larger, it provides more chance for bidirectional search to prune unnecessary valuation. DivMODis has a comparable performance with BiMODis, as it mainly benefits from the bi-directional implementation as BiMODis, hence exploits early pruning and incurs a small overhead in performing a stream-style placement strategy. (2) BiMODis, NOBiMODis and DivMODis is 2.5, 2 and 2.5 times faster than ApxMODis on average, respectively. ApxMODis takes most time in exploring a large universal table with a reduction only strategy. It is not sensitive to ϵ , as its localized search from “data rich” end may be stabilized at a local optimal ϵ -Pareto set that are less frequently updated.

T_1 : Varying maxl. Fixing $\epsilon = 0.2$ and varying maxl from 2 to 6, Fig. 9(b) tells us that all MODis algorithms take longer time as maxl becomes larger. As expected, there are more states with non- ϵ -dominance relation to existing solution to be resolved, and more state nodes to be valued. On the other hand, ApxMODis is the most sensitive to maxl due to rapid growth of search space. BiMODis is much less sensitive, as it mitigates the impact better due to bi-directional strategy and pruning.

T_3 : Varying ϵ and maxl. We report the efficiency of our algorithms for regression task. Our observation is consistent with their counterparts for T_1 . This verifies that the efficiency of our approach is not very sensitive to the type of learning tasks or models.

We have also investigated the impact of input data size. Due to limited space, we report the details in [2].

Exp-3: Case study. We next perform two real-world case studies to illustrate the application scenarios of our MODis methods. We anonymized the cases for this submission.

“Find data with models”. A material science team has a trained random forest-based classifier to recognize peaks in a 2-dimensional angle-intensive X-ray diffraction data. They hope to find good additional datasets for which the model has improved performance

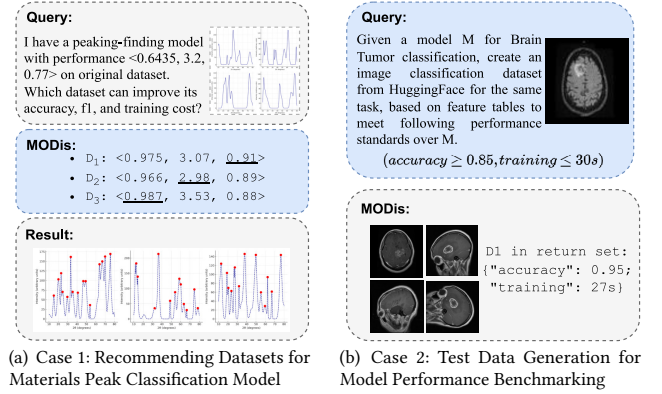


Figure 10: Case Analysis

in terms of F_1 score, training cost, and accuracy for downstream fine-tuning. The team has uploaded their original X-ray datasets and models to a crowd-sourced X-ray data collection platform we have deployed, with a performance vector $\langle 0.6435, 3.2, 0.77 \rangle$. Over a set of shared X-ray diffraction datasets from other facilities, BiMODis created three new datasets $\{D_1, D_2, D_3\}$, which optimizes the model in each measure to 0.975, 3.07, and 0.89, respectively. A test of the model over the datasets is illustrated in Fig. 10, which is manually validated to be accurate with ground-truth given by a database provided by a third-party international institution.

Generating Test data for Model Evaluation. In the second case study, we show that our configurable MODis paradigm readily fits the need for generating data (from scratch) for model benchmarking with required performances [34]. Given a trained scientific image classification model from Kaggle, we processed a pool of image features \mathcal{D} from HF with 75 tables and in total 768 columns and more than 1000 rows as data sources. We set the ranges of training time requirement to be “accuracy > 0.85”, “training cost” <30s, By setting these as *hard constraints* for BiMODis, we found that it outputs a set of 5 datasets within 30 seconds, over which the model’s record on each measure is 0.95 in accuracy and 27s in training costs. This maps to a set of test images that serve as better training or testing data, with an example illustrated in Fig. 10(b).

9 CONCLUSION

We have introduced a novel problem of multi-objective data discovery (MODis). We formalized a class of data discovery system and their computation in terms of finite state transducers, and verified their properties in terms of terminating, confluence and expressiveness. We then introduced three feasible algorithms to approximate Pareto set of datasets with different search strategies, and show that they preserve the provable sub-optimality in terms of ϵ -Pareto set. Our experimental study has verified their feasibility in creating new datasets to improve multiple measurements for model performance. One future topic is to study the incremental maintenance of the datasets upon changes of the source datasets. Another topic is to develop effective parallel algorithms for distributed data discovery.

REFERENCES

- [1] U.S. General Services Administration. 2023. Data.gov. <https://www.data.gov/>
- [2] Anonymous. 2023. Full Version. https://anonymous.4open.science/w/modis_full/full.pdf
- [3] Dan Brickley, Matthew Burgess, and Natasha Noy. 2019. Google Dataset Search: Building a search engine for datasets in an open Web ecosystem. In *The World Wide Web Conference*. 1365–1375.
- [4] Amit Chakrabarti and Sagar Kale. 2015. Submodular maximization meets streaming: matchings, matroids, and more. *Mathematical Programming* (2015).
- [5] AnHai Doan, Alon Halevy, and Zachary Ives. 2012. *Principles of data integration*. Elsevier.
- [6] Marwa El Halabi, Slobodan Mitrović, Ashkan Norouzi-Fard, Jakab Tardos, and Jakub Tarnawski. 2020. Fairness in Streaming Submodular Maximization: Algorithms and Hardness. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*.
- [7] European Organization For Nuclear Research and OpenAIRE. 2013. Zenodo. <https://doi.org/10.25495/7GXX-RD71>
- [8] Grace Fan, Jin Wang, Yuliang Li, Dan Zhang, and Renée Miller. 2022. Semantics-aware Dataset Discovery from Data Lakes with Contextualized Column-based Representation Learning. *arXiv preprint arXiv:2210.01922* (2022).
- [9] Sainyam Galhotra, Yue Gong, and Raul Castro Fernandez. 2023. METAM: Goal-Oriented Data Discovery. *arXiv preprint arXiv:2304.09068* (2023).
- [10] Christian Glaßer, Christian Reitwießner, Heinz Schmitz, and Maximilian Witek. 2010. Approximability and hardness in multi-objective optimization. In *Programs, Proofs, Processes: 6th Conference on Computability in Europe, CiE 2010, Ponta Delgada, Azores, Portugal, June 30–July 4, 2010. Proceedings 6*. 180–189.
- [11] Zhiqiang Gong, Ping Zhong, and Weidong Hu. 2019. Diversity in machine learning. *IEEE Access* 7 (2019), 64323–64350.
- [12] H2O.ai. 2022. *H2O: Scalable Machine Learning Platform*. <https://github.com/h2oai/h2o-3> version 3.42.0.2.
- [13] José Hernández-Orallo, Wout Schellaert, and Fernando Martínez-Plumed. 2022. Training on the test set: Mapping the system-problem space in AI. In *AAAI*, Vol. 36. 12256–12261.
- [14] Hugging Face AI 2023. Hugging Face – The AI Community Building the Future. <https://huggingface.co/>
- [15] Kaggle. 2023. Kaggle: Your Home for Data Science. <https://www.kaggle.com/>
- [16] Aamod Khatiwada, Roei Shraga, Wolfgang Gatterbauer, and Renée J Miller. 2022. Integrating Data Lake Tables. *Proceedings of the VLDB Endowment* 16, 4 (2022), 932–945.
- [17] Ron Kohavi and George H John. 1997. Wrappers for feature subset selection. *Artificial intelligence* 97, 1-2 (1997), 273–324.
- [18] Hsiang-Tsung Kung, Fabrizio Luccio, and Franco P Preparata. 1975. On finding the maxima of a set of vectors. *J. ACM* 22, 4 (1975), 469–476.
- [19] Yan-Nei Law, Haixun Wang, and Carlo Zaniolo. 2004. Query languages and data models for database sequences and data streams. In *VLDB*. 492–503.
- [20] Maurizio Lenzerini. 2002. Data integration: A theoretical perspective. In *PODS*.
- [21] Guoliang Li. 2017. Human-in-the-loop data integration. *Proceedings of the VLDB Endowment* 10, 12 (2017), 2006–2017.
- [22] Jundong Li, Kewei Cheng, Suhang Wang, Fred Morstatter, Robert P Trevino, Jiliang Tang, and Huan Liu. 2017. Feature selection: A data perspective. *ACM computing surveys (CSUR)* 50, 6 (2017), 1–45.
- [23] Yao Lu, Aakanksha Chowdhery, Srikanth Kandula, and Surajit Chaudhuri. 2018. Accelerating machine learning inference with probabilistic predicates. In *Proceedings of the 2018 International Conference on Management of Data*. 1493–1508.
- [24] Jianyu Miao and Lingfeng Niu. 2016. A survey on feature selection. *Procedia computer science* 91 (2016), 919–926.
- [25] Xuan-Phi Nguyen, Shafiq Joty, Kui Wu, and Ai Ti Aw. 2020. Data Diversification: A Simple Strategy For Neural Machine Translation. In *Advances in Neural Information Processing Systems*.
- [26] Xuan Vinh Nguyen, Jeffrey Chan, Simone Romano, and James Bailey. 2014. Effective global approaches for mutual information based feature selection. In *Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery and data mining*. 512–521.
- [27] Andrei Paleyes, Raoul-Gabriel Urma, and Neil D Lawrence. 2022. Challenges in deploying machine learning: a survey of case studies. *Comput. Surveys* 55, 6 (2022), 1–29.
- [28] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. 2011. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* 12 (2011), 2825–2830.
- [29] Hanchuan Peng, Fuhui Long, and Chris Ding. 2005. Feature selection based on mutual information criteria of max-dependency, max-relevance, and min-redundancy. *IEEE Transactions on pattern analysis and machine intelligence* 27, 8 (2005), 1226–1238.
- [30] Yuji Roh, Geon Heo, and Steven Euijong Whang. 2019. A survey on data collection for machine learning: a big data-ai integration perspective. *IEEE Transactions on Knowledge and Data Engineering* 33, 4 (2019), 1328–1347.
- [31] Darius Šidlauskas and Christian S Jensen. 2014. Spatial joins in main memory: Implementation matters! *Proceedings of the VLDB Endowment* 8, 1 (2014), 97–100.
- [32] Tom F Sterkenburg and Peter D Grünwald. 2021. The no-free-lunch theorems of supervised learning. *Synthese* 199, 3-4 (2021), 9979–10015.
- [33] George Tsagouris and Christos Zaroliagis. 2009. Multiobjective optimization: Improved FPTAS for shortest paths and non-linear objectives with applications. *Theory of Computing Systems* 45, 1 (2009), 162–186.
- [34] Margarida LC Vicente, José FO Granjo, Ruomu Tan, and Franz D Böhner. 2022. A Benchmark Model to Generate Batch Process Data for Machine Learning Testing and Comparison. In *Computer Aided Chemical Engineering*, Vol. 51. Elsevier, 217–222.
- [35] Kaichao You, Yong Liu, Jianmin Wang, and Mingsheng Long. 2021. Logme: Practical assessment of pre-trained models for transfer learning. In *International Conference on Machine Learning*. 12133–12143.
- [36] Huaqing Zhang, Jian Wang, Zhanquan Sun, Jacek M Zurada, and Nikhil R Pal. 2019. Feature selection for neural networks using group lasso regularization. *IEEE Transactions on Knowledge and Data Engineering* 32, 4 (2019), 659–673.
- [37] Patrick Ziegler and Klaus R Dittrich. 2007. Data integration—problems, approaches, and perspectives. In *Conceptual modelling in information systems engineering*. Springer, 39–58.